

ECE 6930-004

HPC Fault Tolerance

BASICS OF CHECKPOINT RESTART

DR. JON C. CALHOUN

Schadenfreude! Late October 2012

Peer 1 hosting located in lower Manhattan had backup generators located on the 18th floor

Hurricane Sandy came ashore flooding the building's lobby and basement

Unfortunately, the emergency generator fuel pumping system was located in the basement

Due to post 9/11 restrictions sufficient supplies of fuel were not kept on site

Schadenfreude! Late October 2012

To prevent downtime, a bucket brigade was formed to carry fuel to the 17th floor where the fuel tanks were located

Peer 1 engineers and customers worked tirelessly to fill the fuel tanks

No downtime was reported

Oddly enough bucket brigade was not part of the Peer 1 disaster recovery plan



We NEED volunteers

Date	Paper/Topic	Presenter
8/23	Introduction/Syllabus/What is HPC	Calhoun
8/28	Basic Fault Tolerance Concepts	Calhoun
8/30	Modeling Reliability / Basic Failure Detection	Calhoun
9/4	Toward Exascale Resilience	Calhoun
9/6	Lessons Learned From the Analysis of System Failures at Petascale: The Case of Blue Waters	Omar
9/11	Basics of Checkpoint-restart	Calhoun
9/13	Evaluation of Simple Causal Message Logging for Large-Scale Fault Tolerant HPC Systems	
9/28	Design, Modeling, and Evaluation of a Scalable Multi-level Checkpointing System	
9/20	MCRENGINE: A Scalable Checkpointing System Using Data-Aware Aggregation and Compression	
9/25	What is a soft error?	Calhoun

When an application fails

How to detect failures?

Heart beats



Fail-stop failures often result in the crash of the application

Without some mechanism to save and restart the application during its execution, it must restart from the beginning

Methods of recovery

How to recover from fail-stop failures?

Redundancy

Replication

Checkpointing

Due to resource overheads HPC primarily leverages checkpointing to recover applications

Checkpoint-restart

Let's explore how checkpoint-restart and variations on it are used in HPC

What should we checkpoint?

Checkpointing can be used at many different levels (e.g., instruction, function, application)

Here we will look at checkpointing the application

System level

Serializes the entire process's image

- Memory allocations
- Processor registers
- File handles

Can be done by:

- OS (BLCR)
- Compiler (C3)
- External library (libckpt)



System level checkpointing requires little modifications to source code

- Can use preemption

System level

What is there not to like about system level checkpointing?

Although it requires minimal modifications to the application, system level checkpoint has several key drawbacks:

- Not portable
- Job not malleable
- Size of checkpoint is large



Let's leverage user-level knowledge to address these issues

Application level

The user serializes the state of the process

PROS

Smaller checkpoint sizes

Portability

Custom file formats

May support job malleability

CONS

Difficult to implement if preemption is required

- Random location in call graph
- Keep all the local temporary variables

More work on the programmer

When to checkpoint?

Determining checkpointing time can be based on

- Elapsed time
- Number of messages received or sent count

Synchronous/Blocking:

- All computation stops and all pending communications must finish
- Application is fully serialized before computation and communication can resume

Asynchronous:

- Application does not need to halt when it is time to checkpoint
- System level can use copy-on-write with new process
- Application level saves state around key locations

Where to checkpoint?

The checkpoint must survive the failure of system components; therefore, it must reside in non-volatile storage

The need for true persistence does not preclude the use of volatile memories

- RAM
- Neighbor nodes

The checkpoints taken in volatile memories are periodically written to non-volatile storages

What does checkpointing look like in practice?

To understand what happens when checkpointing, let's look at the interactions of 3 processes though time

On restart, computation halts, rolls back to the previous checkpoint, and computation resumes

- All processes or just failed?

Initially let's look at the case when every process must restart

- Mitigate risks of the [domino effect](#), [missing messages](#), [orphan messages](#)

Checkpointing notation

Each process P_i has local state denoted $LS_{i,k}$ at some time k

Let:

$send(m_{ij})$ be a send event of a message m_{ij} by P_i to P_j

$recv(m_{ij})$ be a receive event of message m_{ij} by P_j

$time(x)$ be the time that event x occurs

Furthermore,

$send(m_{ij}) \in LS_{i,k}$ **iff** $time(send(m_{ij})) < time(LS_{i,k})$

$recv(m_{ij}) \in LS_{j,k}$ **iff** $time(recv(m_{ij})) < time(LS_{j,k})$

Checkpointing notation

For processes P_i and P_j let's define two message sets to denote a message's status

Transit:

$$transit(LS_{i,k}, LS_{i,k}) = \{m_{ij} \mid send(m_{ij}) \in LS_{i,k} \wedge recv(m_{ij}) \notin LS_{j,k}\}$$

- Message is in flight

Inconsistent:

$$inconsistent(LS_{i,k}, LS_{i,k}) = \{m_{ij} \mid send(m_{ij}) \notin LS_{i,k} \wedge recv(m_{ij}) \in LS_{j,k}\}$$

Checkpointing notation

Let the global state of the application (p processes) be defined as

$$GS = \{LS_0, LS_1, \dots, LS_p, \}$$

The global state is said to be **consistent** iff

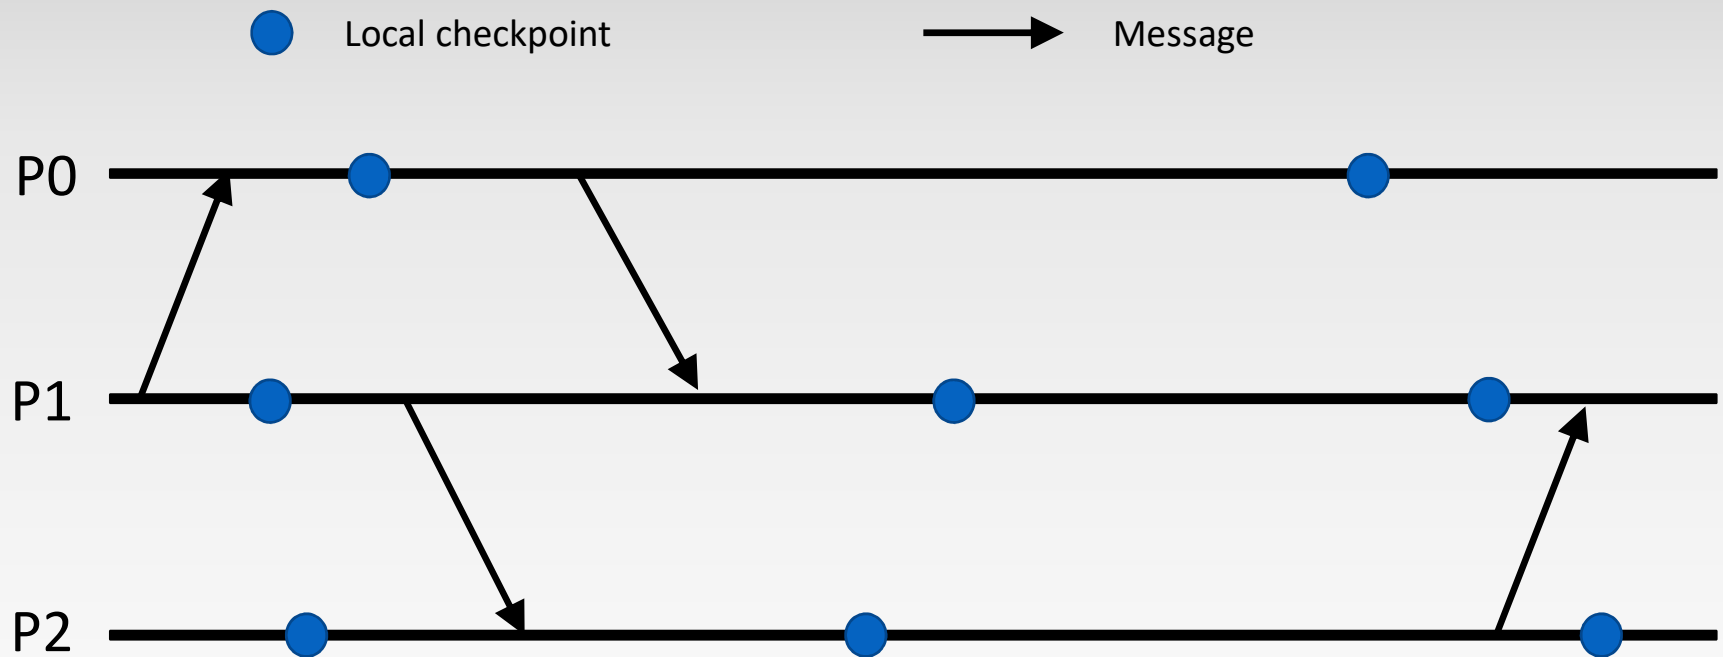
$$\forall i, \forall j : 0 \leq i, j \leq p :: inconsistent(LS_i, LS_j) = \emptyset$$

The global state is said to be **transitless** iff

$$\forall i, \forall j : 0 \leq i, j \leq p :: transit(LS_i, LS_j) = \emptyset$$

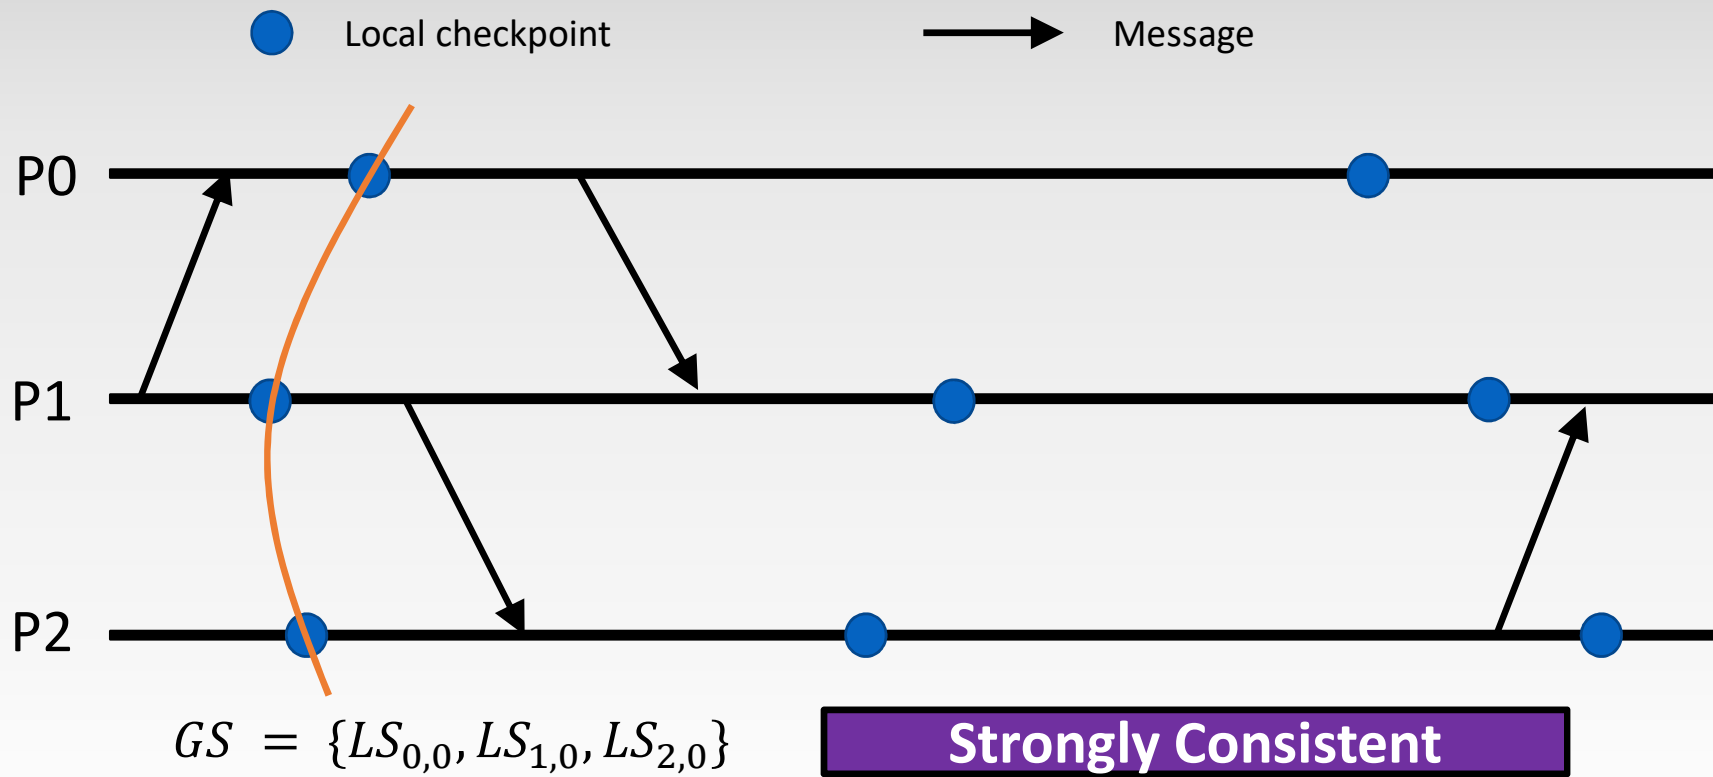
The global state is **strongly consistent** if it is both **consistent** and **transitless**

Example




What **global state** make sense?

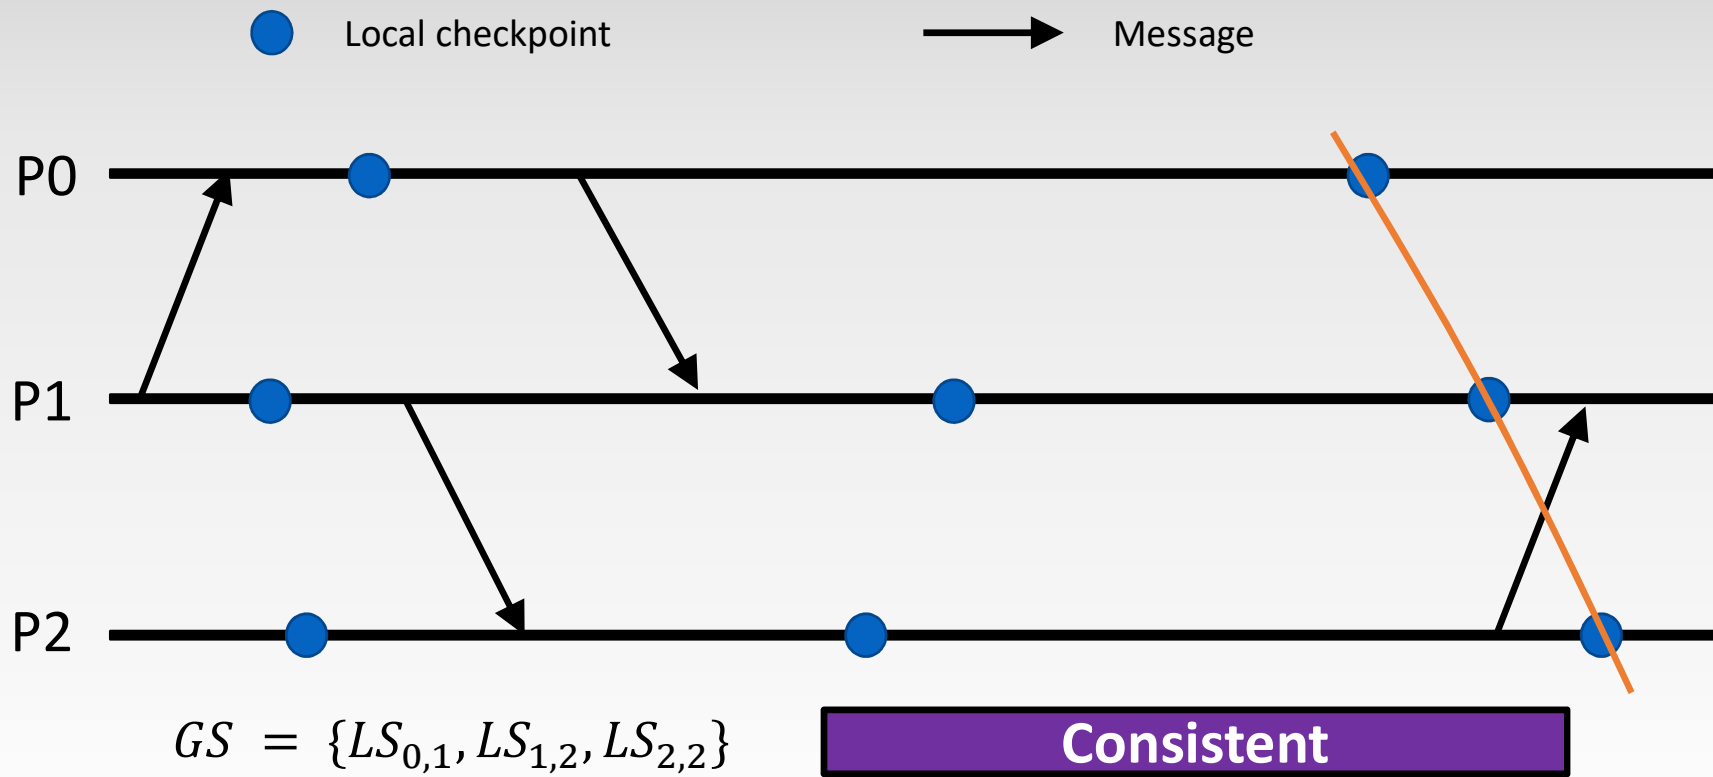
Example



Example



Example



Determining the cut

To have a valid application checkpoint we need at least **consistent cut** of checkpoints

How to make the cut?

Coordinated approach: Use **marker messages** to indicate that a checkpoint is being taken

Uncoordinated approach: Attempt to form a consistent global cut at recovery time

- **Domino effect** must be overcome